

leastsquares

March 22, 2024

```
[74]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import fetch_data
import math
```

```
[2]: ls_data = fetch_data.getStandings()
```

```
[3]: pts_array = ls_data['pts'].values
gd_array = ls_data['gd'].values
```

0.0.1 Least Squares Applications

In statistics, it is common to represent the variable of interest Y as a linear combination of other variables (factors) x_1, \dots, x_n as $Y = \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_n X_n + \varepsilon$ where ε is a random variable that accounts for the missing factors.

Before we begin, let's talk about Least Squares.

In some cases $Ax = b$ may not have a solution, if the system is not compatible ($b \notin \text{img}(A)$)

In such case, we solve the least squares problem. That is, we want the vector $x \in \mathbb{R}^n$ that minimizes $\|Ax - b\|^2$

0.0.2 Geometry of Least Squares

We can think about it like this: Let's say we have a plane for the $\text{img}(A)$ and we make a vector $\hat{w} = A\hat{x}$, this vector would be in the $\text{img}(A)$ because we are multiplying any vector in \mathbb{R}^n times our matrix A . Now since our vector ($b \notin \text{img}(A)$), we need to find a vector that minimizes the distance between our vector w and b . We call this the residual, $\hat{r} = b - A\hat{x}$, the residual is what measures how close $A\hat{x}$ is to b .

$$\left\| \begin{bmatrix} b_1 - w_1 \\ b_2 - w_2 \\ \dots \\ b_n - w_n \end{bmatrix} \right\|^2 = (b_1 - w_1)^2 + (b_2 - w_2)^2 + \dots + (b_n - w_n)^2 = \hat{r}$$

This example is where the least squares terminology comes from, we are trying to find the least squares estimate (least squared distances). Now the closest vector in any subspace to a vector that is not in the subspace is called the projection. So, $\hat{w} = A\hat{x}$ is called the projection of b onto $\text{img}(A)$. Now through the fundamental theorem of linear algebra, our residual \hat{r} lies in the $\text{coker}(A)$ because

the residual is orthogonal to the $\text{img}(A)$. Now we know, the “best” solution is one for which the residual is orthogonal to the $\text{img}(A)$, since the $\text{coker}(A) = (\text{img}(A))^\perp$.

0.0.3 Derivation

Through the information we now know, we can derive the “normal equation”.

Since, $b - A\hat{x} \in \text{coker}(A)$ then $A^\top(b - A\hat{x}) = 0$

$A^\top b - A^\top A\hat{x}$ then translates to $(A^\top A\hat{x} = A^\top b)$

In the case for Linear Regression, the normal equation would then just be $(X^\top X)\hat{\beta} = X^\top y$

Furthermore, solving for β we get $\beta = (X^\top X)^{-1}X^\top y$

0.0.4 Now lets do Linear Regression

```
[30]: #constructing our first vector (Goal Differences)
```

```
x1 = np.array(gd_array).reshape(-1,1)
```

```
#constructing our design matrix
```

```
ones_array = np.ones_like(x1)
```

```
X = np.hstack((ones_array, x1))
```

```
#constructing our column response vector (Points)
```

```
y = np.array(pts_array)
```

```
[49]: print('The design matrix is:')
```

```
print(X[:5]) #showing an example of the first 5 rows of the design matrix
```

```
print('The data vector is:')
```

```
print(y[:5]) #showing an example of the first 5 observations of the data vector
```

The design matrix is:

```
[[ 1 73]
 [ 1 50]
 [ 1 61]
 [ 1 49]
 [ 1 68]]
```

The data vector is:

```
[93 88 89 90 92]
```

Computing the equation $\beta = (X^\top X)^{-1}X^\top y$

```
[50]: beta = np.dot((np.dot(np.linalg.inv(np.dot(X.T, X)), X.T)), y)
```

```
print('The least squares coefficients are \n')
```

```
print(beta)
```

The least squares coefficients are

```
[51.51670337  0.55603074]
```

```
[86]: fig, ax = plt.subplots()


#getting our residuals vector
residual_vector = y - np.dot(X, beta)

ax.plot(x1, y, linestyle='none', markersize=4, marker='o', color='blue')

#plotting our least squares line
ax.plot(x1, beta[0] + beta[1]*x1, 'r')
#plotting residual lines
for i in range(len(x1)):
    ax.plot([x1[i], x1[i]], [y[i], y[i] - residual_vector[i]], 'k--',
            linewidth=0.5)

ax.set_xlabel('Goal Differences (x)')
ax.set_ylabel('Points (y)')
```

```
[86]: Text(0, 0.5, 'Points (y)')
```



0.0.5 Now let's see how many points a team with a goal difference of +48 can expect

```
[55]: print('y = %d points' % math.floor(beta[0] + beta[1]*49))
```

y = 78 points

```
[91]: #calculate the mean of our response vector
mean_y = np.mean(y)
#calculate the total sum of squares (TSS)
TSS = np.sum((y - mean_y)**2)
#calculate R squared
R_squared = 1 - (residual_vector**2).sum() / TSS
print("R squared:", R_squared)
```

R squared: 0.7218755174333396